

## Exercise # 3

---

**Q3a)** For this question we are going to derive both the expectation step and the maximization step independently:

**Expectation Step:** To begin we are given that  $S_k$  is diagonal, which gives us the following properties:

$$|S_k| = \sigma_1^2 \times \sigma_2^2 \times \dots \times \sigma_n^2 = \prod_i^n \sigma_i^2$$

$$S_k^{-1} = \begin{bmatrix} \frac{1}{\sigma_1^2} & \dots & 0 \\ \vdots & & \vdots \\ 0 & & \frac{1}{\sigma_n^2} \end{bmatrix}$$

Continuing from the slides the we get that:

$$\begin{aligned} r_{ik} &= q_i(Z_i = k) \\ &= p_\theta(Z_i = k | x_i) \\ &= \frac{p_\theta(Z_i = j, x_i)}{p_\theta(x_i)} \\ &= \frac{\pi_k N(\mu_k, S_k, x_i)}{\sum_{l=1}^k \pi_k N(\mu_l, S_l, x_i)} \end{aligned}$$

Note that the denominator is calculated outside of the red step, and the red step is only the numerator. Therefore any optimization for a diagonal matrix will have to occur within  $p_\theta(Z_i = j, x_i)$ , thus we get the following expansion:

$$\begin{aligned} \pi_k N(\mu_k, S_k, X_i) &= \frac{1}{\sqrt{|2\pi S_k|}} \exp\left(-\frac{1}{2}(x_i - \mu_k)^T S_k^{-1}(x_i - \mu_k)\right) \\ &= \frac{1}{\sqrt{2\pi^k |S_k|}} \exp\left(-\frac{1}{2}\left(\frac{(x_{i1} - \mu_{k1})^2}{\sigma_1^2} + \frac{(x_{i2} - \mu_{k2})^2}{\sigma_2^2} + \dots + \frac{(x_{in} - \mu_{kn})^2}{\sigma_n^2}\right)\right) \\ &= \frac{\exp\left(-\frac{1}{2}\left(\frac{(x_{i1} - \mu_{k1})^2}{\sigma_1^2} + \frac{(x_{i2} - \mu_{k2})^2}{\sigma_2^2} + \dots + \frac{(x_{in} - \mu_{kn})^2}{\sigma_n^2}\right)\right)}{\sqrt{2\pi^k \sigma_1^2 \times \sigma_2^2 \times \dots \times \sigma_n^2}} \\ &= \frac{\exp\left(-\frac{1}{2}\frac{(x_{i1} - \mu_{k1})^2}{\sigma_1^2}\right)}{\sqrt{2\pi\sigma_1^2}} \times \frac{\exp\left(-\frac{1}{2}\frac{(x_{i2} - \mu_{k2})^2}{\sigma_2^2}\right)}{\sqrt{2\pi\sigma_2^2}} \times \dots \times \frac{\exp\left(-\frac{1}{2}\frac{(x_{in} - \mu_{kn})^2}{\sigma_n^2}\right)}{\sqrt{2\pi\sigma_n^2}} \end{aligned}$$

**Maximization Step** As given in the slides are our is to maximize the following:

$$p_{\theta}(x) = \sum_i^k \pi_k N(\mu_k, S_k, x_i)$$

Which from the slides can we know can be rewritten as:

$$\begin{aligned} &= \operatorname{argmax}_{\theta} \sum_{i=1}^n \sum_{j=1}^k q_i(Z_i = j) \log p_{\theta}(x_i, Z_i = j) \\ &= \operatorname{argmax}_{\theta} \sum_{i=1}^n \sum_{j=1}^k q_i(Z_i = j) \log \left[ \frac{\pi_k}{\sqrt{|2\pi S_k|}} \exp \left( -\frac{1}{2} (x_i - \mu_k)^T S_k^{-1} (x_i - \mu_k) \right) \right] \\ &= \operatorname{argmax}_{\theta} \sum_{i=1}^n \sum_{j=1}^k q_i(Z_i = j) \left[ \log(\pi_k) - \frac{k}{2} \log(2\pi) - \frac{1}{2} \log(|S_k|) + \left( -\frac{1}{2} (x_i - \mu_k)^T S_k^{-1} (x_i - \mu_k) \right) \right] \\ &= \operatorname{argmax}_{\theta} \sum_{i=1}^n \sum_{j=1}^k r_{ik} \left[ \log(\pi_k) - \frac{k}{2} \log(2\pi) - \frac{1}{2} \log(|S_k|) + \left( -\frac{1}{2} (x_i - \mu_k)^T S_k^{-1} (x_i - \mu_k) \right) \right] \end{aligned}$$

We can then take the derivative w.r.t to  $S_k$  and set to zero, to find where this is concave function is maximized:

$$\begin{aligned} \sum_{i=1}^n \sum_{j=1}^k r_{ik} \left[ \log(\pi_k)' - \frac{k}{2} \log(2\pi)' - \frac{1}{2} \log(|S_k|)' + \left( -\frac{1}{2} (x_i - \mu_k)^T S_k^{-1} (x_i - \mu_k) \right)' \right] &= 0 \\ \sum_{i=1}^n \sum_{j=1}^k r_{ik} \left[ 0 - 0 - \frac{1}{2} \log \left( \prod_i^n \sigma_i^2 \right)' + \left( -\frac{1}{2} (x_i - \mu_k)^T (x_i - \mu_k) \right) \right] &= 0 \\ \sum_{i=1}^n \sum_{j=1}^k r_{ik} \left[ 0 - 0 - \frac{1}{\sigma_{ik}^2} + \left( -\frac{1}{2} (x_i - \mu_k)^T (x_i - \mu_k) \right) \right] &= 0 \end{aligned}$$

We can rearrange to get:

$$\begin{aligned} \sum_{i=1}^n \frac{r_{ik}}{\sigma_{ik}^2} &= \sum_{i=1}^n \frac{1}{2} (x_i - \mu_k)^T (x_i - \mu_k) \\ S_k &= \sum_{i=1}^n \frac{\frac{1}{2} (x_i - \mu_k)^T (x_i - \mu_k)}{r_{ik}} \end{aligned}$$

On the next page is my implementation of the algorithm.

---

```

import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
from numpy.random import RandomState

TOL = 1e-5

def estep(x_i, mu, s, d):
    currentSum = 0
    for idx in range(0, d):
        step1 = (x_i[idx] - mu[idx]) ** 2 / s[idx] * -0.5
        step2 = np.log(np.sqrt(2 * np.pi * np.abs(s[idx])))
        currentSum += step1 - step2
    return currentSum

def mstep(x, r, elementSum, mu, cluster, d, N):
    step1 = 0
    for i in range(N):
        step1 += np.exp(r[i, cluster]) * x.iloc[i, d] ** 2

    step2 = step1 / elementSum[cluster]
    return step2 - mu[cluster, d] ** 2

def train(X, maxIter, k):
    N, D = X.shape

    pi = np.ones(k) / k
    mu = np.random.rand(k, D)

    r = np.zeros((N, k))
    s = np.ones((k, D))

    pastError = 0

    log_likelihood = 0

    for iter in range(maxIter):

        print("Start of E Step")
        # E step
        clusterSum = np.zeros(N)
        for i in range(N):
            for cluster in range(k):
                r[i, cluster] = np.log(pi[cluster])
                r[i, cluster] += estep(X.iloc[i], mu[cluster], s[cluster], D)

            for cluster in range(k):
                clusterSum[i] += np.exp(r[i, cluster])

            for cluster in range(k):
                r[i, cluster] -= np.log(clusterSum[i])

        # Error Check
        pastError = log_likelihood
        log_likelihood = 0
        for i in range(N):
            step3 = 0

```

```

        for cluster in range(k):
            step1 = X.iloc[i] - mu[cluster]
            step2 = np.exp(-0.5 * step1 @ np.linalg.inv(np.diag(s[cluster])) @ step1.T)
            step3 += step2 / np.sqrt(pi[cluster] * np.prod(s[cluster]))
            log_likelihood += np.log(step3)

log_likelihood = log_likelihood * -1

if iter > 1 and pastError - log_likelihood <= TOL * log_likelihood:
    break
print("Current Error is: ", log_likelihood)
# M Step
elementSum = np.zeros(k)
for cluster in range(k):
    for i in range(N):
        elementSum[cluster] += np.exp(r[i, cluster])

for cluster in range(k):
    pi[cluster] = elementSum[cluster] / N

for cluster in range(k):
    for d in range(D):
        localSum = 0
        for i in range(N):
            localSum += (np.exp(r[i, cluster])) * X.iloc[i, d]
        mu[cluster, d] = localSum / elementSum[cluster]

for cluster in range(k):
    for d in range(D):
        s[cluster, d] = mstep(X, r, elementSum, mu, cluster, d, N)

print("Done M Step")
return log_likelihood, pi, mu, r, s

gmm = pd.read_csv("gmm_dataset.csv", header=None)

yVals = []
xVals = []

last = 0
for i in range(1, 11):
    log_likelihood, pi, mu, r, s = train(gmm, 500, i)

    print(i)
    xVals.append(i)
    yVals.append(log_likelihood)
    last = log_likelihood

plt.plot(xVals, yVals)
plt.xlabel("Number of Gaussian Distributions (k)")
plt.ylabel("Negative Log Likelihood")
plt.title("")
plt.show()

```

---

**Time and Space Complexity Analysis:** For the expectation step we need to loop through each cluster and element of  $x$ , as well as each feature which gives us a time complexity of  $O(NDK)$ . For the loss function we need to only loop through clusters and  $N$  so we get a time complexity of  $O(ND)$ . For the maximization step we need to loop through each element of  $x$ , each feature and each cluster so its only  $O(NDK)$  time. Therefore the total time complexity is:

$$O(NDK)$$

As for space complexity we need to store the following variables:

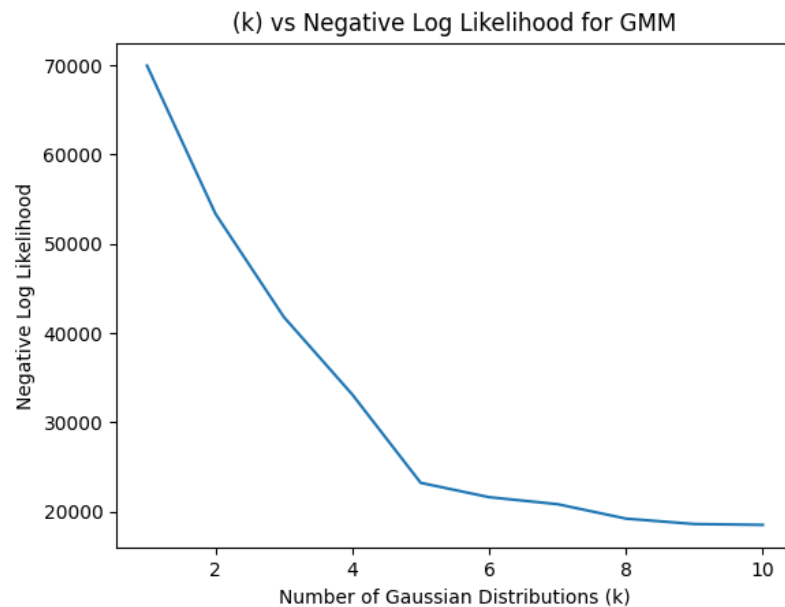
$$pi = O(K)$$

$$mu = O(KD)$$

$$r = O(NK)$$


$$s = O(KD)$$

**Graph and Analysis:**



The best choice of  $k$  in this case is 5. As we can see in the graph that's where a hinge occurs so after that point the marginal benefit in log loss decreases (but is not negative). **The model parameters are provided in a separate csvs.**

**Q3b)** For this question we get the following graph, where for  $k = 5$  we got a test error of roughly 22.17% seen below is the graph:



The code used to generate this graph is seen below, although repeated members aren't included:

---

```
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import tensorflow as tf

from sklearn.decomposition import PCA

def normalize_img(image):
    scaled_image = image / 255. # Scale the image to 0-1

    return scaled_image.flatten()

(x_train, y_train), (x_test, y_test) = tf.keras.datasets.mnist.load_data()

dataToTrain = []
dataToTest = []

for i in range(0, len(x_train)):
    dataToTrain.append(normalize_img(x_train[i, :]))

for i in range(0, len(x_test)):
    dataToTest.append(normalize_img(x_test[i, :]))

pca = PCA(n_components=100)
pca.fit(dataToTrain)

dataToTrain = pca.transform(dataToTrain)
```

```

dataToTest = pca.transform(dataToTest)

dataToTrainByDigit= [[] for i in range(10)]

for i in range(0, len(x_train)):
    dataToTrainByDigit[y_train[i]].append(dataToTrain[i,:])

predictionPerDigit = []
for i in range(0, 10):
    predictionPerDigit.append(len(dataToTrainByDigit[i])/len(x_train))

xVals = []
yVals = []
for k in range(1, 6):
    correct = 0
    incorrect = 0

    modelsPerDigit = []

    for i in range(0, 10):
        modelsPerDigit.append(trainGMM(k, dataToTrainByDigit[i]))

    for test_idx in range(0, len(dataToTest)):

        prob = calculatePredictions(modelsPerDigit[0], dataToTest[test_idx])
        digitWithHighestProb = 0

        for i in range(1, 10):

            likelihood = calculatePredictions(modelsPerDigit[i], dataToTest[test_idx])
            if prob < likelihood * predictionPerDigit[i]:
                prob = likelihood * predictionPerDigit[i]
                digitWithHighestProb = i

        if digitWithHighestProb == y_test[test_idx]:
            correct += 1
        else:
            incorrect += 1

    xVals.append(k)
    yVals.append(incorrect/ (correct+incorrect))
    print(k, ": ", incorrect/ (correct+incorrect))

plt.plot(xVals, yVals)
plt.xlabel("Number of Gaussian Distributions (k)")
plt.ylabel("Test Error")
plt.title("(k) vs Test Error")
plt.show()

```

---