

# CS 480 Cheat Sheet

## Perceptron

We assume  $(x_1, y_1), \dots, (x_n, y_n)$  belongs to some distribution. Choose predictive function  $h_n$  such that  $\max Pr(h(x_i) = y_i)$

**Dot Product:**  $\langle w, x \rangle = \sum w_i x_i$

**Padding:**  $\langle w, x \rangle + b = \langle (w, b), (x, 1) \rangle$

Note:  $z = (w, b)$ ,  $a_i = y_i \langle x_i, 1 \rangle$

**Linear Seperable:** if  $s > 0$  and  $Az > s1$

If data is not linearly seperable perceptron stalls.

Margin is determined by closest point to the hyperplane.

Perceptron finds a solution, may not be best solution.

**l2 Norm:**  $\|x\|_2 = \sqrt{\sum_i x_i^2}$

**Error Bound**  $\leq \frac{R^2 \|z\|_2^2}{s^2}$ ,  $R = \max \|a_i\|_2$

**Margin:**  $\gamma = \max_{\|z\|_2=1} \min_i \langle a_i, z \rangle$

**One-versus-all:**  $\hat{y} = \operatorname{argmax}_k w_k^T x + b_k$

**One-versus-one:**  $\#\{x^T w_{k,k'} + b_{k,k'} > 0, x^T w_{k',k} + b_{k',k} < 0\}$

## Linear Regression

**Gradient:** if  $f(x) \mathbb{R}^d \rightarrow \mathbb{R}$ ,  $\Delta f(v) = \left( \frac{\delta f}{\delta v_1}, \dots, \frac{\delta f}{\delta v_d} \right) \mathbb{R}^d \rightarrow \mathbb{R}^d$

**Hessian:**  $\Delta^2 f(v) = \begin{bmatrix} \frac{\delta^2 f}{\delta v_1^2} & \dots & \delta v_d^2 \delta v_1^2 \\ \vdots & & \vdots \\ \frac{\delta^2 f}{\delta v_1^2 \delta v_d^2} & \dots & \delta^2 v_d^2 \end{bmatrix} : \mathbb{R}^d \rightarrow \mathbb{R}^{d \times d}$

**Empirical Risk Minimization:**  $\operatorname{argmin}_w \frac{1}{n} \sum_d l_w(x, y)$

**Convexity #1:**  $f(\lambda x_1 + (1-\lambda)x_2) \leq \lambda f(x_1) + (1-\lambda)f(x_2)$

**Convexity #2:** if  $\Delta^2 f(x)$  is positive semi definite.

**Positive Semidefinite:** if  $M \in \mathbb{R}^{d \times d}$ , PSD iff  $v^T M v \geq 0$

Loss function needs to be convex to optimize.

Setting loss function to 0 optimizes our solution.

MLE principle: pick parameters that maximize likelihood.

**Ridge Regularization:**  $\operatorname{argmin}_w \|Aw - z\|_2^2 + \lambda \|w\|_2^2$

**Lasso Regularization:**  $\operatorname{argmin}_w \|Aw - z\|_2^2 + \lambda \|w\|_2$

## k-Nearest Neighbour Classification

**Bays Optimal Classifier:**  $f^*(x) = \operatorname{argmax}_c Pr(y = c|x)$

**NN Assumption:**  $Pr[y = c|x] \approx Pr[y' = c|x'], x \approx x'$

No classifier can do as good as bayes.

Can't be described as a parameter vector.

Can express non linear relationships.

Takes 0 training time, and  $O(nd)$  to  $O(d \log n)$  testing time.

Small values of k lead to overfitting.

Large values of k lead to high error.

**1NN Limit:** as  $n \rightarrow \infty$  then  $L_{1NN} \leq 2L_{Bayes}(1 - L_{Bayes})$

## Logistic Regression

Classifications are taken into account confidence:  $\hat{y} \in (-1, 1)$

**Bernoulli Model:**  $Pr[y = 1|x, w] = p(x, w) \in (-1, 1)$

**Logit Transform:**  $\log \left( \frac{p(x, w)}{1-p(x, w)} \right) = \langle x, w \rangle = \frac{1}{1+\exp(-\langle x, w \rangle)}$

**Optimizing Loss:**  $\Delta_w l_w(x_i, y_i) = (p_i(x_i, w) - y_i)x_i$

**Iterative Update:**  $w_t = w_{t-1} - \eta d_t$

**Gradient Descent:**  $d_t = \frac{1}{n} \sum_i \Delta_w l_{wt-1}(x_i, y_i)$

**Stochastic GD:**  $B \in [n], d - t = \frac{1}{|B|} \sum_{i \in B} \Delta_w l_{wt-1}(x_i, y_i)$

**Newton's Method**  $d_t$  is given by the equation below:

$$d_t = \left( \frac{1}{n} \sum_i \Delta_w^2 l_{wt-1}(x_i, y_i) \right)^{-1} \left( \frac{1}{n} \sum_i \Delta_w l_{wt-1}(x_i, y_i) \right)$$

**Multiclass Logistic Regression** where  $k = \text{class}$ :

$$Pr[y = k|x, w] = \frac{\exp(\langle w_k, x \rangle)}{\sum_i \exp(\langle w_i, x \rangle)}$$

## Hard-Margin SVM

Assume that dataset is linearly separable. Hard Margin SVM's will try to find the "best" solution. The best solution is the one that maximizes margin.

**Optimize:**  $\min_{w, b} \frac{1}{2} \|w\|_2^2 \text{ s.t. } y\hat{y} \geq 1$

**Primal:**  $\min_{w, b} \frac{1}{2} \|w\|_2^2 \text{ s.t. } y_i \langle w, x_i \rangle + b \geq 1$

**Dual:**  $\min_{\alpha} \frac{1}{2} \sum_i \sum_j \alpha_i \alpha_j y_i y_j \langle x_i, x_j \rangle \text{ s.t. } \sum \alpha_i y_i = 0$

**Complimentary Slackness:**  $a_i (y_i (\langle w, x_i \rangle + b) - 1) = 0, \forall i$

**Support Vector:** if  $a_i > 0$  then  $w = \sum a_i y_i x_i$

## Soft-Margin SVM

Data does not need to be linearly separable.

**Soft-Margin:**  $\min_{w, b} \frac{1}{2} \|w\|_2^2 + C \sum_i \max(0, 1 - y_i \hat{y}_i)$

if  $1 - y_i \hat{y}_i \leq 0 \implies$  Correct side of margin.

if  $0 < 1 - y_i \hat{y}_i \leq 1 \implies$  Correctly classified, inside of margin.

if  $y_i \hat{y}_i \leq 0 \implies$  incorrectly classified.

If  $C=0$  ignore data, if  $C=\infty$ , hard-margin.

**Slack Variable:** define  $\gamma_i$  such that  $\max(0, 1 - y_i \hat{y}_i) \leq \gamma_i$

**Split in Two:**  $0 \leq \gamma_i$  and  $1 - y_i \hat{y}_i \leq \gamma_i$

**Dual Solution:** Note  $0 \leq \gamma_i$  and  $1 - y_i \hat{y}_i \leq \gamma_i$  implies:

$$= \max_{\alpha, \beta} \min_{w, b, \gamma} \frac{1}{2} \|w\|_2^2 + \sum (C \gamma_i + \alpha (1 - y_i \hat{y}_i - \gamma_i) - \beta_i \gamma_i)$$

$$= \min_{\alpha} \frac{1}{2} \sum_i \sum_j \alpha_i \alpha_j y_i y_j \langle x_i, x_j \rangle - \sum a_i \text{ s.t. } \sum a_i y_i = 0$$

if  $a_i = 0$  then  $y_i = 0$ , point is classified correctly.

if  $a_i > 0$  and  $y_i = 0$ , point is on margin.

if  $a_i > 0$  and  $y_i > 0$ , point is on within margin.

**Loss Function:**  $L = \frac{C}{n} \sum_i l_{w, b}(x_i, y_i) + \frac{1}{2} \|w\|_2^2$

**Gradient Descent:**  $\frac{\delta L}{\delta w} = w + C/N \sum \delta_i$

if  $1 - y_i \hat{y}_i \geq 0$ , then  $\delta = -y_i x_i$  else  $\delta = 0$

## Kernels

Map data to new space where it is linearly separable.

**Padding Trick:**  $\phi(x) = [w, 1]$  and  $w = \langle x, p \rangle$

**New Classifier:**  $\langle \phi(x), w \rangle = \langle x, p \rangle + b > 0$

**Quadratic Feature:**  $x^T Q x + \sqrt{2} x^T p + b$ , which gives us:

$$\phi(x) = [x^T, \sqrt{2}x, 1] \text{ and } w = [Q, p, b]$$

With feature map  $\phi: \mathbb{R}^d \rightarrow \mathbb{R}^{d \times d + d + 1}$ , time  $O(d)$  to  $O(d^2)$

This can take infinite time in high dimensions. For the dual we only need to calculate dot product.

**Kernel:**  $k: \mathbb{R}^d \times \mathbb{R}^d \rightarrow \mathbb{R}$  if  $k(x, x') = l \angle \phi(x), \phi(x')$

A kernel is valid if it's symmetric and positive semi-definite.

**New Kernel:**  $K_{ij} = \langle \phi(x^i), \phi(x^j) \rangle = k(x^i, x^j)$

**Classify New:**  $\operatorname{sign}(\sum a^i y^i k(x^i, x))$

**Polynomial Kernel t:**  $k(\langle x, x' \rangle + 1)^t$

**Gaussian Basis:**  $\exp(-\|x - x'\|_2^2)$

SVM (Linear Kernel):  $O(nd)$  train time,  $O(d)$  test time.

General Kernel:  $O(n^2 d)$  train time  $O(nd)$  test time.

## Decision Trees

Can do classification or regression, handle non-linear functions.

May fail on linear functions.

Start at one node, and split each. Select the pure node.

**Loss:**  $t^* = \operatorname{argmin}_t l(\{(x_i, y_i) : x_i \leq t\}) + l(\{(x_i, y_i) : x_i > t\})$

Let  $p_c = \text{frac of } S \text{ with label } c$ .  $\hat{y} = \operatorname{argmax}_c p_c$

**Misclassification Loss:**  $l(s) = 1 - p_y$

**Entropy Loss:**  $l(s) = - \sum_{\text{classes } c} p_c \log p_c$

**Gini Index Loss:**  $l(s) = \sum_{\text{classes } c} p_c (1 - p_c)$

**Regression:**  $l(s) = \min_p \sum_{i \in S} (y_i - p)^2 = \sum_{i \in S} (y_i - \bar{y}_s)^2$

We can stop based on run time, depth or splits.

Once a tree is fully grown we can prune it.

## Bagging

Training on empirical mean gives a variance of:  $E[\hat{\mu}] = \mu$

$$Var[\hat{\mu}] = Var[\frac{1}{n} \sum X_i] = \frac{1}{n^2} Var[\sum X_i] = \sigma^2/n$$

We can reduce variance by taking a sample of B points:

$$Var[\hat{\mu}] = Var[\frac{1}{B} \sum X_i] = \frac{1}{B^2} Var[\sum X_i] = \sigma^2/Bn$$

We can sample these points with replacement, and in practice this will work.

We aggregate by doing regression  $f(x) = \frac{1}{B} \sum f^j(x)$ .

Classification done by majority vote.

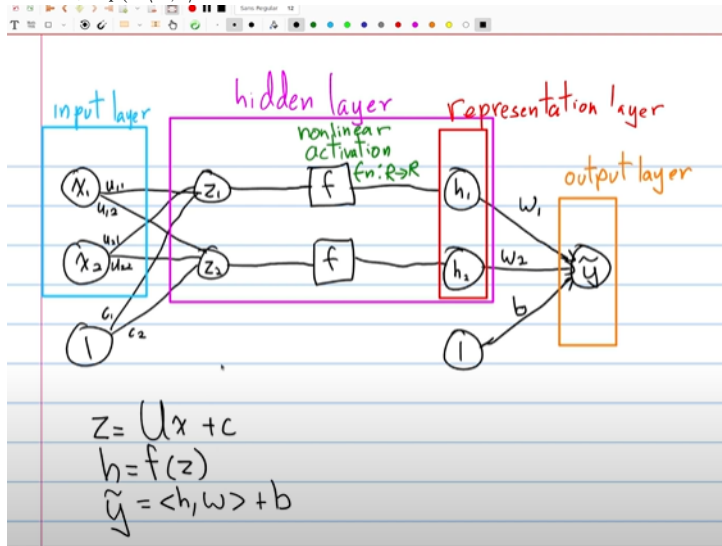
**Random Forests** Bootstrap but select  $\sqrt{d}$  features.

## Multilayer Perceptron

Neural Networks learn mapping from the data.

**Sigmoid:**  $\sigma(t) = \frac{1}{1+e^{-t}}$

$$\hat{y} = \frac{1}{1+\exp(-(w,x)-b)}$$



**ReLU(t)** =  $\max(0, t)$

**Loss**  $l_\theta(x, y) = -\sum y_i \log y_i$

**Tanh(t)** =  $\frac{e^t - e^{-t}}{e^t + e^{-t}}$

**Gradient Descent**  $\theta^t = \theta^{t-1} - \eta \Delta L_{\theta^{t-1}}$

**Chain rule:**  $\frac{dz}{dx} = \frac{dz}{dy} \frac{dy}{dx}$

Any continuous function can be approximated well by a 2 layer nn.

## Deep Networks

Most neural networks have many parameters. We can minimize overfitting with:

Regularization loss, gradient descent, and equivalently:

$$\theta_t \leftarrow (1 - \eta\lambda)\theta_{t-1} - \eta\Delta L_{\theta_{t-1}}(x, y)$$

We can drop nodes and use normalization of features:

mean =  $\frac{1}{n} \sum X_i$ ,  $X_i \leftarrow X_i - \mu$ ,  $\sigma_j^2 = \frac{1}{n} \sum X_{i,j}^2$ ,  $X_{i,j} = X_{i,j}/\sigma_j$

We do normalization on each batch, such that:

$$z^i = W^i h^{i-1} + b^i \text{ or } h^i = f(z^i)$$

We can do normalization on each neuron (batchnorm) or each layer.

**Batch GD:**  $\theta \leftarrow \theta - \eta * \frac{1}{n} \sum \delta l_\theta(x_i, y_i)$ , optimize gradient.

**Momentum:**  $v_t = \gamma v_{t-1} + (1 - \gamma)\mu$ ,  $\theta_t \leftarrow \theta_{t-1} - v_t$  **RMSPprop** let  $g \in R^p$ ,  $G_{t,i} = \sum^t g_{t,i}^2$  and  $\theta_t \leftarrow \theta_{t-1} - \frac{\mu}{\sqrt{G_{t,i} + \epsilon}} g_{t,i}$

- $\beta_1, \beta_2, \epsilon$  hyperparameters
  - $\beta_1 = 0.9, \beta_2 = 0.999, \epsilon = 10^{-8}$
- $m_{t,i} = \beta_1 m_{t-1,i} + (1 - \beta_1) g_{t,i}$  (momentum)
- $v_{t,i} = \beta_2 v_{t-1,i} + (1 - \beta_2) g_{t,i}^2$  (RMSprop)
- $\hat{m}_{t,i} = \frac{m_{t,i}}{1 - \beta_1^t}$ ,  $\hat{v}_{t,i} = \frac{v_{t,i}}{1 - \beta_2^t}$
- $\theta_{t,i} \leftarrow \theta_{t-1,i} - \frac{\eta}{\sqrt{\hat{v}_{t,i} + \epsilon}} \hat{m}_{t,i}$

## [Perceptron] Algorithm

**Algorithm:** The Perceptron (Rosenblatt 1958)

**Input:** Dataset  $\mathcal{D} = \{(x_i, y_i) \in \mathbb{R}^d \times \{\pm 1\} : i = 1, \dots, n\}$ , initialization  $w \in \mathbb{R}^d$  and  $b \in \mathbb{R}$ , threshold  $\delta \geq 0$

**Output:** approximate solution  $w$  and  $b$

```

1 for t = 1, 2, ... do
2   receive training example index  $I_t \in \{1, \dots, n\}$  // the index  $I_t$  can be random
3   if  $y_{I_t}(w^T x_{I_t} + b) \leq \delta$  then
4      $w \leftarrow w + y_{I_t} x_{I_t}$  // update only after making a "mistake"
5      $b \leftarrow b + y_{I_t}$ 

```

## [Perceptron] Error Bound

• **Theorem (informal):** If  $\exists z, s$  such that  $Az \geq s\mathbf{1}$ , perceptron makes at most  $R^2 \|z\|_2^2 / s^2$  mistakes, where  $R = \max \|a_i\|_2$ .

• Pick the "best" one to minimize  $\|z\|_2^2 / s^2$  and thus the number of mistakes

$$\min_{(z,s): Az \geq s\mathbf{1}} \frac{\|z\|_2^2}{s^2} = \min_{(z,s): \|z\|_2=1, Az \geq s\mathbf{1}} \frac{1}{s^2}$$

$$= \frac{1}{(\max_{\|z\|_2=1} \min_{Az \geq s\mathbf{1}} \langle a_i, z \rangle)^2} = \frac{1}{\gamma^2}$$

$\gamma = \max_{\|z\|_2=1} \min_i \langle a_i, z \rangle$  is the margin of the solution wrt the dataset.

## [Linear Regression] Convexity of Loss Function

• Loss fn  $\|Aw - z\|_2^2 = (Aw - z)^T (Aw - z) = (w^T A^T - z^T)(Aw - z)$   
 $= w^T A^T Aw - z^T Aw - w^T A^T z + z^T z$   
 $= w^T A^T Aw - 2w^T A^T z + z^T z$

• Claim: if  $f(x) = x^T Ax + x^T b + c$ , then  $\nabla f(x) = (A + A^T)x + b$

• Thus  $\nabla_w \|Aw - z\|_2^2 = 2A^T Aw - 2A^T z$

• Checking the Hessian,  $\nabla_w^2 \|Aw - z\|_2^2 = 2A^T A \geq 0$

• Why? Since  $2v^T A^T A v = 2\|Av\|_2^2 \geq 0$  for any vector  $v$

## [Linear Regression] Deriving MLE

$y = \langle w, x \rangle + z$ , where  $z \sim N(0, \sigma^2)$

$$\hat{w} = \arg \max_w \Pr(x_1, y_1, \dots, x_n, y_n | w)$$

$$= \arg \max_w \prod_i \Pr(x_i, y_i | w)$$

$$= \arg \max_w \prod_i \Pr[y_i | x_i, w] \Pr[x_i | w]$$

$$= \arg \max_w \prod_i \Pr[y_i | x_i, w]$$

$$= \arg \max_w \prod_i \Pr[y_i | x_i, w]$$

$$= \arg \max_w \log \left( \prod_i \Pr[y_i | x_i, w] \right)$$

$$= \arg \max_w \log \left( \prod_i \Pr[y_i | x_i, w] \right)$$

$$= \arg \max_w \sum_i \log(\Pr[y_i | x_i, w])$$

$$\text{(Note: } y_i | x_i, w \sim N(\langle w, x_i \rangle, \sigma^2) \text{)}$$

$$= \arg \max_w \sum_i \log \left( \frac{1}{\sqrt{2\pi\sigma^2}} \exp \left( -\frac{(y_i - \langle w, x_i \rangle)^2}{2\sigma^2} \right) \right)$$

$$= \arg \max_w \sum_i \log \left( \frac{1}{\sqrt{2\pi\sigma^2}} \right) + \log \left( \exp \left( -\frac{(y_i - \langle w, x_i \rangle)^2}{2\sigma^2} \right) \right)$$

$$= \arg \max_w \sum_i -\frac{(y_i - \langle w, x_i \rangle)^2}{2\sigma^2}$$

$$= \arg \min_w \sum_i (y_i - \langle w, x_i \rangle)^2$$

$$= \arg \min_w \sum_i (y_i - \langle w, x_i \rangle)^2$$

Loss function is the squared error!

## [Linear Regression] Cross Validation

• Split training data into  $k$  sets (draw on board), e.g.  $k = 10$  is common  
 For each  $\lambda$ :

For  $i = 1$  to  $k$ :

$w_{\lambda,i}$  = train on all data but split  $i$  with hyperparameter  $\lambda$

$\text{perf}_{\lambda,i}$  = performance of  $w_{\lambda,i}$  on the split  $i$

$\text{perf}_\lambda = \sum_i \text{perf}_{\lambda,i}$

Return  $\lambda$  which has the biggest  $\text{perf}_\lambda$

## [KNN] Algorithm

•  $\Pr_{y \sim D_{Y|X=x}}[y = c|x] \approx \Pr_{y' \sim D_{Y|X=x'}}[y' = c|x']$  when  $x$  and  $x'$  are close

**Algorithm:** kNN

**Input:** Dataset  $\mathcal{D} = \{(x_i, y_i) \in X \times Y : i = 1, \dots, n\}$ , new instance  $x \in X$ , hyperparameter  $k$

**Output:**  $y = y(x)$

1 for  $i = 1, 2, \dots, n$  do

2  $d_i \leftarrow \text{dist}(x, x_i)$

// avoid for-loop if possible

3 find indices  $i_1, \dots, i_k$  of the  $k$  smallest entries in  $d$

4  $y \leftarrow \text{aggregate}(y_{i_1}, \dots, y_{i_k})$

## [Logistic Regression] MLE of $\hat{w}$

$$\hat{w} = \arg \min_w \frac{1}{n} \sum_{i=1}^n \log(\exp(-y_i \langle x_i, w \rangle) + \exp((1 - y_i) \langle x_i, w \rangle))$$

Let  $\tilde{y}_i = +1$  if  $y_i = 1$ , and  $\tilde{y}_i = -1$  if  $y_i = 0$ .

$$\hat{w} = \arg \min_w \frac{1}{n} \sum_{i=1}^n \log(1 + \exp(-\tilde{y}_i \langle x_i, w \rangle))$$

## [Logistic Regression] Logit Transform

$$\log\left(\frac{p(x, w)}{1 - p(x, w)}\right) = \langle x, w \rangle$$

$$\frac{p(x, w)}{1 - p(x, w)} = \exp(\langle x, w \rangle) \text{ (LHS: "odds ratio")}$$

$$p(x, w) = \exp(\langle x, w \rangle) (1 - p(x, w))$$

$$p(x, w) = \exp(\langle x, w \rangle) - \exp(\langle x, w \rangle) p(x, w)$$

$$p(x, w)(1 + \exp(\langle x, w \rangle)) = \exp(\langle x, w \rangle)$$

$$p(x, w) = \frac{\exp(\langle x, w \rangle)}{1 + \exp(\langle x, w \rangle)}$$

$$p(x, w) = \frac{1}{1 + \exp(-\langle x, w \rangle)} \triangleq \text{sigmoid}(\langle x, w \rangle)$$

## [Logistic Regression] Deriving MLE #1

$$\hat{w} = \arg \max_w \prod_{i=1}^n \Pr[(x_i, y_i) | w]$$

$$= \arg \max_w \prod_{i=1}^n p(x_i, w)^{y_i} (1 - p(x_i, w))^{1-y_i}$$

(Let  $p_i = p(x_i, w)$ )

$$= \arg \max_w \log \prod_{i=1}^n p_i^{y_i} (1 - p_i)^{1-y_i}$$

$$= \arg \max_w \sum_{i=1}^n \log(p_i^{y_i} (1 - p_i)^{1-y_i})$$

$$= \arg \max_w \sum_{i=1}^n y_i \log p_i + (1 - y_i) \log(1 - p_i)$$

("cross entropy loss")

Recall:  $p(x, w) = \frac{1}{1 + \exp(-\langle x, w \rangle)}$

Suppose some  $y_i = 1$ . Then

$$y_i \log p_i + (1 - y_i) \log(1 - p_i) = \log p_i$$

$$= \log(1 + \exp(-\langle x_i, w \rangle))^{-1}$$

$$= -\log(1 + \exp(-\langle x_i, w \rangle))$$

Similarly, if  $y_i = 0$ , then

$$y_i \log p_i + (1 - y_i) \log(1 - p_i) = \log(1 - p_i)$$

$$= \log\left(\frac{\exp(-\langle x_i, w \rangle)}{1 + \exp(-\langle x_i, w \rangle)}\right)$$

$$= -\log(1 + \exp(-\langle x_i, w \rangle))$$

## [Logistic Regression] Deriving MLE #2

$$\hat{w} = \arg \max_w \sum_{i=1}^n y_i \log p_i + (1 - y_i) \log(1 - p_i)$$

If  $y_i = 1$ , argument is  $-\log(1 + \exp(-\langle x_i, w \rangle))$

If  $y_i = 0$ , argument is  $-\log(1 + \exp(\langle x_i, w \rangle))$

$$\hat{w} = \arg \max_w \sum_{i=1}^n -\log(\exp(-y_i \langle x_i, w \rangle) + \exp((1 - y_i) \langle x_i, w \rangle))$$

$$\hat{w} = \arg \min_w \frac{1}{n} \sum_{i=1}^n \log(\exp(-y_i \langle x_i, w \rangle) + \exp((1 - y_i) \langle x_i, w \rangle))$$

## [Logistic Regression] $d_t$ Selection

- Gradient Descent
  - $d_t = \frac{1}{n} \sum_{i=1}^n \nabla_w \ell_{w_{t-1}}(x_i, y_i)$  (note, = 0 at optimum)
  - Running time?
- Stochastic gradient descent
  - Draw random set  $B \subseteq [n]$ , then let  $d_t = \frac{1}{|B|} \sum_{i \in B} \nabla_w \ell_{w_{t-1}}(x_i, y_i)$
- Newton's Method
  - $d_t = \left(\frac{1}{n} \sum_{i=1}^n \nabla_w^2 \ell_{w_{t-1}}(x_i, y_i)\right)^{-1} \left(\frac{1}{n} \sum_{i=1}^n \nabla_w \ell_{w_{t-1}}(x_i, y_i)\right)$
  - Often needs fewer steps to converge, but more time/memory per step

## [Hard Margin SVM] Goal

- SVM: Like perceptron, but try to maximize margin
 
$$\max_{w', b'} \gamma, \text{ s.t. } \|w'\|_2 = 1, y_i(\langle w', x_i \rangle + b') \geq \gamma \text{ for all } i$$
- Substitute  $w' = \gamma w, b' = \gamma b$ 

$$\max_{\gamma, w, b} \gamma, \text{ s.t. } \|w\|_2 = 1/\gamma, y_i(\langle w, x_i \rangle + b) \geq \gamma \text{ for all } i$$

$$\max_{\gamma, w, b} \gamma, \text{ s.t. } \|w\|_2 = 1/\gamma, y_i(\langle w, x_i \rangle + b) \geq 1 \text{ for all } i$$

$$\max_{\gamma, w, b} \frac{1}{\|w\|_2}, \text{ s.t. } y_i(\langle w, x_i \rangle + b) \geq 1 \text{ for all } i$$

$$\min_{w, b} \frac{1}{2} \|w\|_2^2 \text{ s.t. } y_i(\langle w, x_i \rangle + b) \geq 1 \text{ for all } i$$

## [Hard Margin SVM] Dual Formation

$$\max_{\alpha \in \mathbb{R}^n, \alpha \geq 0} \min_{w, b} \frac{1}{2} \|w\|_2^2 - \sum_{i=1}^n \alpha_i (y_i(\langle w, x_i \rangle + b) - 1)$$

Fix some  $\alpha$  for now, solve inner minimization. How? Set gradient = 0!

$$\frac{\partial}{\partial b} = -\sum_i \alpha_i y_i = 0, \quad \frac{\partial}{\partial w} = w - \sum_i \alpha_i y_i x_i = 0$$

Substitute into above and rearrange...

$$\min_{\alpha \in \mathbb{R}^n, \alpha \geq 0} \frac{1}{2} \sum_i \sum_j \alpha_i \alpha_j y_i y_j \langle x_i, x_j \rangle - \sum_i \alpha_i \text{ s.t. } \sum_i \alpha_i y_i = 0$$

## [Soft Margin SVM] Goal

$$\min_{w, b} \frac{1}{2} \|w\|_2^2 + C \sum_i \max(0, 1 - y_i \hat{y}_i)$$

- (draw cases of  $y_i \hat{y}_i$ , versus hard-margin case)
  - If  $1 - y_i \hat{y}_i \leq 0$ , then on the correct side of margin
  - If  $0 \leq y_i \hat{y}_i \leq 1$ , correctly classified but within margin
  - If  $y_i \hat{y}_i \leq 0$ , incorrectly classified
- (draw hinge loss versus 0-1 loss, perceptron)
- If  $C = 0$ , ignore data, if  $C = \infty$ , hard-margin SVM

## [Soft Margin SVM] Dual Formation

$$\min_{w, b, \gamma} \frac{1}{2} \|w\|_2^2 + C \sum_i \gamma_i \text{ s.t. } 0 \leq \gamma_i \text{ and } 1 - y_i \hat{y}_i \leq \gamma_i \text{ for all } i$$

- Introduce dual variables and take Lagrangian
 
$$\max_{\alpha, \beta \in \mathbb{R}^n, \alpha, \beta \geq 0} \min_{w, b, \gamma} \frac{1}{2} \|w\|_2^2 + \sum_i (C \gamma_i + \alpha_i (1 - y_i \hat{y}_i - \gamma_i) - \beta_i \gamma_i)$$
- Take derivative of inner problem, set to 0, substitute, simplify... (exercise)
 
$$\min_{\alpha \in \mathbb{R}^n, C \geq \alpha \geq 0} \frac{1}{2} \sum_i \sum_j \alpha_i \alpha_j y_i y_j \langle x_i, x_j \rangle - \sum_i \alpha_i \text{ s.t. } \sum_i \alpha_i y_i = 0$$

## [Soft Margin SVM] Optimization

- $\ell_{w, b}(x, y) = \max(0, 1 - y(\langle w, x \rangle + b))$
- Optimize loss function
 
$$L = \frac{C}{n} \sum_i \ell_{w, b}(x_i, y_i) + \frac{1}{2} \|w\|_2^2$$
  - Normalize by  $n$  this time, same problem just rescaled
  - Note similarity to ridge regression ( $C$  on former vs  $\lambda$  on latter)
- Gradient descent:  $\frac{\partial L}{\partial w} = w + \frac{C}{n} \sum_i \delta_i$ 
  - $\delta_i = -y_i x_i$  if  $1 - y_i \hat{y}_i \geq 0$ ,  $\delta_i = 0$  if  $1 - y_i \hat{y}_i \leq 0$  (draw, note non-diff pt)

## [Decision Tree] Example

$$S_L = \{(x_i, y_i) : x_{ij} \leq t\}, S_R = \{(x_i, y_i) : x_{ij} > t\}$$

$$(j^*, t^*) = \arg \min_{j, t} |S_L| \ell(S_L) + |S_R| \ell(S_R)$$

Gini index:  $\sum_{\text{classes } c} \hat{p}_c (1 - \hat{p}_c)$

Split on smokes?

$$\text{No: } \hat{p}_0 = \frac{1}{4}, \hat{p}_1 = \frac{3}{4}; \text{ Yes: } \hat{p}_0 = \frac{2}{3}, \hat{p}_1 = \frac{1}{3}.$$

$$\text{Cost: } 4 \cdot \left(\frac{3}{4}\right) \left(\frac{1}{4}\right) + \left(\frac{1}{4}\right) \left(\frac{3}{4}\right) + 6 \cdot \left(\frac{2}{3}\right) \left(\frac{1}{3}\right) + \left(\frac{1}{3}\right) \left(\frac{2}{3}\right) = 4.16$$

Split on age? (Cheat to save time: use 35 as split)

$$\leq 35: \hat{p}_0 = 1, \hat{p}_1 = 0. > 35: \hat{p}_0 = \frac{1}{6}, \hat{p}_1 = \frac{5}{6}.$$

$$\text{Cost: } 4 \cdot ((0)(1) + (1)(0)) + 6 \cdot \left(\frac{1}{6}\right) \left(\frac{5}{6}\right) + \left(\frac{5}{6}\right) \left(\frac{1}{6}\right) = 1.66$$

Age	Smokes	Cancer?
10	No	0
18	Yes	0
25	No	0
35	Yes	0
50	No	1
55	Yes	1
70	Yes	1
80	No	0
85	Yes	1
90	Yes	1

## [Boosting] Hedge Algorithm

Hedge( $\beta$ ), where  $\beta \in [0,1]$

1. Initialize  $w^{(1)} = [1/n, \dots, 1/n] \in \mathbb{R}^n$

2. For  $t = 1, \dots, T$

- Set  $p^{(t)} = \frac{w^{(t)}}{\sum_i w_i^{(t)}}$  (normalize  $w$  into a distribution)
- Receive loss  $\langle p^{(t)}, \ell^{(t)} \rangle$
- Update  $w_i^{(t+1)} = w_i^{(t)} \beta^{\ell_i^{(t)}}$  (downweight experts based on loss)

Guarantee:  $\sum_{t=1}^T \langle p^{(t)}, \ell^{(t)} \rangle \leq \frac{1}{1-\beta} (\log n + \log(1/\beta) \min_i \sum_{t=1}^T \ell_i^{(t)})$

## [Boosting] AdaBoost Algorithm

1. Initialize  $w^{(1)} = [1/n, \dots, 1/n] \in \mathbb{R}^n$

2. For  $t = 1, \dots, T$

- Set  $p^{(t)} = \frac{w^{(t)}}{\sum_i w_i^{(t)}}$  (normalize  $w$  into a distribution)
- Run WeakLearn on training set (with weights  $p^{(t)}$ )
  - Obtain classifier  $h^{(t)}$  which maps  $(x, y)$  datapoints to  $[0,1]$  (confidence in classification)
- Calculate error  $\epsilon_t = \sum_i p_i^{(t)} |h^{(t)}(x_i) - y_i|$  (should be  $< 0.5$  by WeakLearn guarantees)
- Define  $\beta_t = \frac{\epsilon_t}{1-\epsilon_t}$ , if  $\epsilon_t \leq 1/2$  set  $w_i^{(t+1)} = w_i^{(t)} \beta_t^{1-h^{(t)}(x_i)-y_i}$ 
  - Note: If  $\epsilon_t$  big, then  $\beta_t$  is big. Many errors, so don't downweight points!

3.  $h(x) = 1$  if  $\sum_{t=1}^T \log\left(\frac{1}{\beta_t}\right) h^{(t)}(x) \geq \frac{1}{2} \sum_{t=1}^T \log\left(\frac{1}{\beta_t}\right)$ , 0 else

## [Multilevel Perceptron] 2 Layer Perceptron

- $z = Ux + c, h = f(z), \tilde{y} = \langle h, w \rangle + b$
- Consider:  $U = \begin{bmatrix} 1 & 1 \\ 1 & 1 \end{bmatrix}, c = \begin{bmatrix} 0 \\ -1 \end{bmatrix}, w = \begin{bmatrix} 2 \\ -4 \end{bmatrix}, b = -1$ 
  - Parameters to be learned, but suppose they're just given for now
- Choose  $f(t) = \max(0, t) = \text{ReLU}(t)$  (draw)
  - Activation function – this is a hyperparameter choice
- $x_1 = \begin{bmatrix} 0 \\ 0 \end{bmatrix}, y_1 = -1. z = \begin{bmatrix} 0 \\ -1 \end{bmatrix}, h = \begin{bmatrix} 0 \\ 0 \end{bmatrix}, \tilde{y} = -1$
- $x_2 = \begin{bmatrix} 0 \\ 1 \end{bmatrix}, y_1 = 1. z = \begin{bmatrix} 1 \\ 0 \end{bmatrix}, h = \begin{bmatrix} 1 \\ 0 \end{bmatrix}, \tilde{y} = 2 - 1 = 1$ , etc.

## [Multilevel Perceptron] 3 Layer Perceptron

- (Draw wider three-layer MLP, input  $x \in \mathbb{R}^d$ , output  $\tilde{y} \in \mathbb{R}^m$ )
  - (Illustrate depth and width, representation layer)
- $z^{(1)} = W^{(1)}x, h^{(1)} = f(z^{(1)}), z^{(2)} = W^{(2)}h^{(1)}, \dots$
- What to do with output  $\tilde{y} \in \mathbb{R}^m$ ?
  - Put through *softmax* to get distribution over  $m$  classes (confidences of each)
  - $\hat{y}_i = \frac{\exp(\tilde{y}_i)}{\sum_{j=1}^m \exp(\tilde{y}_j)}$
- What loss function? Use the *cross-entropy* loss
  - $\ell_\theta(x, y) = -\sum_{i=1}^m y_i \log \hat{y}_i$ 
    - Use "one-hot encoding" of  $y$ : if  $y = c$ , then  $y_c = 1$ , and  $y_i = 0$  for other entries
    - $\hat{y} = g_\theta(x)$ , where  $g_\theta$  is a (somewhat complicated) function


## [Multilevel Perceptron] Back Propagation

- Simple case:  $x \in \mathbb{R}, f, g: \mathbb{R} \rightarrow \mathbb{R}$
- Say  $y = g(x), z = f(y) = f(g(x))$
- Chain rule:  $\frac{dz}{dx} = \frac{dz}{dy} \cdot \frac{dy}{dx}$
- More complex:  $z = ux, h = f(z), y = wh, L = g(y)$  (draw)
- Can compute several derivatives easily:  $\frac{dL}{dy}, \frac{dy}{dw}, \frac{dy}{dh}, \frac{dh}{dz}, \frac{dz}{du}$
- But we care about derivatives of  $L$  wrt parameters  $u, w$
- $\frac{dL}{dw} = \frac{dL}{dy} \cdot \frac{dy}{dw}$  and  $\frac{dL}{du} = \frac{dL}{dy} \cdot \frac{dy}{dh} \cdot \frac{dh}{dz} \cdot \frac{dz}{du}$

## [Deep Networks] Momentum

### Momentum

- Keep memory of previous gradient step
- Let  $\gamma < 1$  (say = 0.9)
- $v_t = \gamma v_{t-1} + (1-\gamma)\eta \cdot \frac{1}{|B|} \sum_{i \in B} \nabla_{\theta_{t-1}} \ell_{\theta_{t-1}}(x_i, y_i)$ 
  - New step: weighted sum of old step and current gradient
- $\theta_t \leftarrow \theta_{t-1} - v_t$
- $v_t = 0.1 g_t + 0.1 \cdot 0.9 g_{t-1} + 0.1 \cdot 0.9^2 g_{t-2} + \dots$ 
  - Total coefficient  $1 - \gamma^t$
- Variant: Nesterov momentum



## [A #1] Ridge Regression

$$\begin{aligned} \min_{w \in \mathbb{R}^d, b \in \mathbb{R}} \frac{1}{2n} \left\| \begin{bmatrix} X & \mathbf{1}_n \\ \sqrt{2\lambda n} I_d & 0_d \end{bmatrix} \begin{bmatrix} w \\ b \end{bmatrix} - \begin{bmatrix} y \\ 0_d \end{bmatrix} \right\|_2^2 \\ &= \min_{w \in \mathbb{R}^d, b \in \mathbb{R}} \frac{1}{2n} \left\| \begin{bmatrix} Xw + b\mathbf{1}_n \\ \sqrt{2\lambda n} w \end{bmatrix} - \begin{bmatrix} y \\ 0_d \end{bmatrix} \right\|_2^2 \\ &= \min_{w \in \mathbb{R}^d, b \in \mathbb{R}} \frac{1}{2n} \left\| \begin{bmatrix} Xw + b\mathbf{1}_n - y \\ \sqrt{2\lambda n} w \end{bmatrix} \right\|_2^2 \\ &= \min_{w \in \mathbb{R}^d, b \in \mathbb{R}} \frac{1}{2n} \left[ (Xw + b\mathbf{1}_n - y)^T (\sqrt{2\lambda n} w I_d)^T \begin{bmatrix} Xw + b\mathbf{1}_n - y \\ \sqrt{2\lambda n} w \end{bmatrix} \right] \\ &= \min_{w \in \mathbb{R}^d, b \in \mathbb{R}} \frac{1}{2n} \left[ \|Xw + b\mathbf{1}_n - y\|_2^2 + 2\lambda n \|w\|_2^2 \right] \\ &= \min_{w \in \mathbb{R}^d, b \in \mathbb{R}} \frac{1}{2n} \|Xw + b\mathbf{1}_n - y\|_2^2 + \lambda \|w\|_2^2 \end{aligned}$$

## [A #1] Ridge Regression Derivatives

$$\begin{aligned} &= \frac{1}{2n} \{ [(X^T X) + (X^T X)^T] w + 2X^T (b\mathbf{1}_n - y) \} + 2\lambda w \\ &= \frac{1}{2n} \{ [(X^T X) + (X^T X)] w + 2X^T (b\mathbf{1}_n - y) \} + 2\lambda w \\ &= \frac{1}{2n} \{ [2(X^T X)] w + 2X^T (b\mathbf{1}_n - y) \} + 2\lambda w \\ &= \frac{1}{n} \{ X^T X w + X^T (b\mathbf{1}_n - y) \} + 2\lambda w \\ \frac{\partial}{\partial w} &= \frac{1}{n} X^T (Xw + b\mathbf{1}_n - y) + 2\lambda w \\ \frac{\partial}{\partial b} &= \left[ \frac{1}{2n} \|Xw + b\mathbf{1}_n - y\|_2^2 + \lambda \|w\|_2^2 \right] \\ &= \frac{\partial}{\partial b} \{ \frac{1}{2n} [w^T X^T X w + 2(b\mathbf{1}_n - y)^T X w + (b\mathbf{1}_n - y)^T (b\mathbf{1}_n - y)] + \lambda w^T w \} \\ &= \frac{1}{n} [X^T X w + b\mathbf{1}_n - y] \\ &= \frac{1}{n} X^T [Xw + b\mathbf{1}_n - y] \\ \frac{\partial}{\partial b} &= \frac{1}{n} X^T (Xw + b\mathbf{1}_n - y) \end{aligned}$$

## [A #2] Kernels

$$\begin{aligned} k(x, y) &= \exp(-\alpha(x-y)^2) \\ &= \exp(-\alpha(x^2 + y^2 - 2xy)) \\ &= \exp(-\alpha(x^2 + y^2)) \exp(2\alpha xy) \\ &= \exp(-\alpha(x^2 + y^2)) \left( 1 + \frac{2\alpha xy}{1!} + \frac{(2\alpha xy)^2}{2!} + \frac{(2\alpha xy)^3}{3!} + \dots \right) \\ &= \exp(-\alpha(x^2 + y^2)) \left( 1 + \sqrt{\frac{2\alpha}{1!}} x \cdot \sqrt{\frac{2\alpha}{1!}} y + \sqrt{\frac{(2\alpha)^2}{2!}} x^2 \cdot \sqrt{\frac{(2\alpha)^2}{2!}} y^2 + \sqrt{\frac{(2\alpha)^3}{3!}} x^3 \cdot \sqrt{\frac{(2\alpha)^3}{3!}} y^3 + \dots \right) \\ &= [\exp(-\alpha x^2) \cdot \left( \sum_{n=0}^{\infty} \sqrt{\frac{(2\alpha)^n}{n!}} x^n \right)] \cdot [\exp(-\alpha y^2) \cdot \left( \sum_{n=0}^{\infty} \sqrt{\frac{(2\alpha)^n}{n!}} y^n \right)] \end{aligned}$$

So, the corresponding feature map is  $\phi(t) = \exp(-\alpha t^2) [1, \sqrt{\frac{2\alpha}{1!}} t, \sqrt{\frac{(2\alpha)^2}{2!}} t^2, \dots]^T$ .

Dual representation would be preferred. The primal representation of soft-margin support vector machine would be

$$\min_{w, b} \frac{1}{2} \|w\|_2^2 + C \sum_{i=1}^n \max(0, 1 - y_i(\langle w, x_i \rangle + b))$$

## [A #2] Duels

minimize the squared norm of the difference between the two feature vectors

Ans:

$$\begin{aligned} \min_{w \in \mathbb{R}^d, b \in \mathbb{R}} \frac{1}{2} \|w\|_2^2 + C \sum_{i=1}^n \max\{|y_i - (w^T x_i + b)| - \epsilon, 0\} \\ &= \min_{w \in \mathbb{R}^d, b \in \mathbb{R}} \frac{1}{2} \|w\|_2^2 + C \sum_{i=1}^n \max\{y_i - (w^T x_i + b), (w^T x_i + b) - y_i, 0\} \end{aligned}$$

Let  $w^T x_i + b = \hat{y}_i$

$$\begin{aligned} &= \min_{w \in \mathbb{R}^d, b \in \mathbb{R}} \frac{1}{2} \|w\|_2^2 + C \sum_{i=1}^n \gamma_i \quad \text{s.t. } \forall i \quad \max(y_i - \hat{y}_i, \hat{y}_i - y_i, 0) \leq \gamma_i \\ &= \min_{w \in \mathbb{R}^d, b \in \mathbb{R}} \frac{1}{2} \|w\|_2^2 + C \sum_{i=1}^n \gamma_i \quad \text{s.t. } \forall i \quad 0 \leq \gamma_i, y_i - \hat{y}_i \leq \gamma_i, \hat{y}_i - y_i \leq \gamma_i \\ &= \min_{w \in \mathbb{R}^d, b \in \mathbb{R}} \max_{\alpha, \beta, \theta \geq 0} \frac{1}{2} \|w\|_2^2 + \sum_{i=1}^n [C\gamma_i - \alpha_i \gamma_i + \beta_i (y_i - \hat{y}_i - \epsilon - \gamma_i) + \theta_i (\hat{y}_i - y_i - \epsilon - \gamma_i)] \\ &= \max_{\alpha, \beta, \theta \geq 0} \min_{w \in \mathbb{R}^d, b \in \mathbb{R}} \frac{1}{2} \|w\|_2^2 + \sum_{i=1}^n [C\gamma_i - \alpha_i \gamma_i + \beta_i (y_i - \hat{y}_i - \epsilon - \gamma_i) + \theta_i (\hat{y}_i - y_i - \epsilon - \gamma_i)] \\ \frac{\partial}{\partial w} &= w - \sum_{i=1}^n (\beta_i x_i - \theta_i x_i) = 0, \quad w = \sum_{i=1}^n (\beta_i x_i - \theta_i x_i) \\ \frac{\partial}{\partial b} &= \sum_{i=1}^n -\beta_i + \theta_i = 0, \quad \sum_{i=1}^n \beta_i = \sum_{i=1}^n \theta_i \\ \frac{\partial}{\partial \gamma} &= \sum_{i=1}^n C - \beta_i - \theta_i - \alpha_i = 0, \quad C = \beta_i - \theta_i - \alpha_i \\ &= \max_{\alpha, \beta, \theta \geq 0} \frac{1}{2} \|w\|_2^2 + \sum_{i=1}^n [C\gamma_i + \beta_i y_i - \beta_i \hat{y}_i - \beta_i \epsilon - \beta_i \gamma_i + \theta_i \hat{y}_i - \theta_i y_i - \theta_i \epsilon - \theta_i \gamma_i - \alpha_i \gamma_i] \end{aligned}$$