

# CS 480 Cheat Sheet

## Perceptron

We assume  $(x_1, y_1), \dots, (x_n, y_n)$  belongs to some distribution. Choose predictive function  $h_n$  such that  $\max Pr(h(x_i) = y_i)$

**Dot Product:**  $\langle w, x \rangle = \sum w_i x_i$

**Padding:**  $\langle w, x \rangle + b = \langle (w, b), (x, 1) \rangle$

Note:  $z = (w, b)$ ,  $a_i = y_i \langle x_i, 1 \rangle$

**Linear Seperable:** if  $s > 0$  and  $Az > s1$

If data is not linearly seperable perceptron stalls.

Margin is determined by closest point to the hyperplane.

Perceptron finds a solution, may not be best solution.

**l2 Norm:**  $\|x\|_2 = \sqrt{\sum_i x_i^2}$

**Error Bound**  $\leq \frac{R^2 \|z\|_2^2}{s^2}$ ,  $R = \max \|a_i\|_2$

**Margin:**  $\gamma = \max_{\|z\|_2=1} \min_i \langle a_i, z \rangle$

**One-versus-all:**  $\hat{y} = \operatorname{argmax}_k w_k^T x + b_k$

**One-versus-one:**  $\#\{x^T w_{k,k'} + b_{k,k'} > 0, x^T w_{k',k} + b_{k',k} < 0\}$

## Linear Regression

**Gradient:** if  $f(x) \mathbb{R}^d \rightarrow \mathbb{R}$ ,  $\Delta f(v) = \left( \frac{\delta f}{\delta v_1}, \dots, \frac{\delta f}{\delta v_d} \right) \mathbb{R}^d \rightarrow \mathbb{R}^d$

**Hessian:**  $\Delta^2 f(v) = \begin{bmatrix} \frac{\delta^2 f}{\delta v_1^2} & \dots & \delta v_d^2 \delta v_1^2 \\ \vdots & & \vdots \\ \frac{\delta^2 f}{\delta v_1^2 \delta v_d^2} & \dots & \delta^2 v_d^2 \end{bmatrix} : \mathbb{R}^d \rightarrow \mathbb{R}^{d \times d}$

**Empirical Risk Minimization:**  $\operatorname{argmin}_w \frac{1}{n} \sum_d l_w(x, y)$

**Convexity #1:**  $f(\lambda x_1 + (1-\lambda)x_2) \leq \lambda f(x_1) + (1-\lambda)f(x_2)$

**Convexity #2:** if  $\Delta^2 f(x)$  is positive semi definite.

**Positive Semidefinite:** if  $M \in \mathbb{R}^{d \times d}$ , PSD iff  $v^T M v \geq 0$

Loss function needs to be convex to optimize.

Setting loss function to 0 optimizes our solution.

MLE principle: pick parameters that maximize likelihood.

**Ridge Regularization:**  $\operatorname{argmin}_w \|Aw - z\|_2^2 + \lambda \|w\|_2^2$

**Lasso Regularization:**  $\operatorname{argmin}_w \|Aw - z\|_2^2 + \lambda \|w\|_2$

## k-Nearest Neighbour Classification

**Bays Optimal Classifier:**  $f^*(x) = \operatorname{argmax}_c Pr(y = c|x)$

**NN Assumption:**  $Pr[y = c|x] \approx Pr[y' = c|x'], x \approx x'$

No classifier can do as good as bayes.

Can't be described as a parameter vector.

Can express non linear relationships.

Takes 0 training time, and  $O(nd)$  to  $O(d \log n)$  testing time.

Small values of k lead to overfitting.

Large values of k lead to high error.

**1NN Limit:** as  $n \rightarrow \infty$  then  $L_{1NN} \leq 2L_{Bayes}(1 - L_{Bayes})$

## Logistic Regression

Classifications are taken into account confidence:  $\hat{y} \in (-1, 1)$

**Bernoulli Model:**  $Pr[y = 1|x, w] = p(x, w) \in (-1, 1)$

**Logit Transform:**  $\log \left( \frac{p(x, w)}{1-p(x, w)} \right) = \langle x, w \rangle = \frac{1}{1+\exp(-\langle x, w \rangle)}$

**Optimizing Loss:**  $\Delta_w l_w(x_i, y_i) = (p_i(x_i, w) - y_i)x_i$

**Iterative Update:**  $w_t = w_{t-1} - \eta d_t$

**Gradient Descent:**  $d_t = \frac{1}{n} \sum_i \Delta_w l_{wt-1}(x_i, y_i)$

**Stochastic GD:**  $B \in [n], d - t = \frac{1}{|B|} \sum_{i \in B} \Delta_w l_{wt-1}(x_i, y_i)$

**Newton's Method**  $d_t$  is given by the equation below:

$$d_t = \left( \frac{1}{n} \sum_i \Delta_w^2 l_{wt-1}(x_i, y_i) \right)^{-1} \left( \frac{1}{n} \sum_i \Delta_w l_{wt-1}(x_i, y_i) \right)$$

**Multiclass Logistic Regression** where  $k = \text{class}$ :

$$Pr[y = k|x, w] = \frac{\exp(\langle w_k, x \rangle)}{\sum_i \exp(\langle w_i, x \rangle)}$$

## Hard-Margin SVM

Assume that dataset is linearly separable. Hard Margin SVM's will try to find the "best" solution. The best solution is the one that maximizes margin.

**Optimize:**  $\min_{w,b} \frac{1}{2} \|w\|_2^2 \text{ s.t. } y\hat{y} \geq 1$

**Primal:**  $\min_{w,b} \frac{1}{2} \|w\|_2^2 \text{ s.t. } y_i \langle w, x_i \rangle + b \geq 1$

**Dual:**  $\min_a \frac{1}{2} \sum_i \sum_j a_i a_j y_i y_j \langle x_i, x_j \rangle \text{ s.t. } \sum a_i y_i = 0$

**Complimentary Slackness:**  $a_i (y_i (\langle w, x_i \rangle + b) - 1) = 0, \forall i$

**Support Vector:** if  $a_i > 0$  then  $w = \sum a_i y_i x_i$

## Soft-Margin SVM

Data does not need to be linearly separable.

**Soft-Margin:**  $\min_{w,b} \frac{1}{2} \|w\|_2^2 + C \sum_i \max(0, 1 - y_i \hat{y}_i)$

if  $1 - y_i \hat{y}_i \leq 0 \implies$  Correct side of margin.

if  $0 < 1 - y_i \hat{y}_i \leq 1 \implies$  Correctly classified, inside of margin.

if  $y_i \hat{y}_i \leq 0 \implies$  incorrectly classified.

If  $C=0$  ignore data, if  $C=\infty$ , hard-margin.

**Slack Variable:** define  $\gamma_i$  such that  $\max(0, 1 - y_i \hat{y}_i) \leq \gamma_i$

**Split in Two:**  $0 \leq \gamma_i$  and  $1 - y_i \hat{y}_i \leq \gamma_i$

**Dual Solution:** Note  $0 \leq \gamma_i$  and  $1 - y_i \hat{y}_i \leq \gamma_i$  implies:

$$= \max_{\alpha, \beta} \min_{w, b, \gamma} \frac{1}{2} \|w\|_2^2 + \sum (C \gamma_i + \alpha (1 - y_i \hat{y}_i - \gamma_i) - \beta_i \gamma_i)$$

$$= \min_{\alpha} \frac{1}{2} \sum_i \sum_j \alpha_i \alpha_j y_i y_j \langle x_i, x_j \rangle - \sum a_i \text{ s.t. } \sum a_i y_i = 0$$

if  $a_i = 0$  then  $y_i = 0$ , point is classified correctly.

if  $a_i > 0$  and  $y_i = 0$ , point is on margin.

if  $a_i > 0$  and  $y_i > 0$ , point is on within margin.

**Loss Function:**  $L = \frac{C}{n} \sum_i l_{w,b}(x_i, y_i) + \frac{1}{2} \|w\|_2^2$

**Gradient Descent:**  $\frac{\delta L}{\delta w} = w + C/N \sum \delta_i$

if  $1 - y_i \hat{y}_i \geq 0$ , then  $\delta = -y_i x_i$  else  $\delta = 0$

## Kernels

Map data to new space where it is linearly separable.

**Padding Trick:**  $\phi(x) = [w, 1]$  and  $w = \langle x, p \rangle$

**New Classifier:**  $\langle \phi(x), w \rangle = \langle x, p \rangle + b > 0$

**Quadratic Feature:**  $x^T Q x + \sqrt{2} x^T p + b$ , which gives us:

$$\phi(x) = [x^T, \sqrt{2}x, 1] \text{ and } w = [Q, p, b]$$

With feature map  $\phi: \mathbb{R}^d \rightarrow \mathbb{R}^{d \times d + d + 1}$ , time  $O(d)$  to  $O(d^2)$

This can take infinite time in high dimensions. For the dual we only need to calculate dot product.

**Kernel:**  $k: \mathbb{R}^d \times \mathbb{R}^d \rightarrow \mathbb{R}$  if  $k(x, x') = l \angle \phi(x), \phi(x')$

A kernel is valid if it's symmetric and positive semi-definite.

**New Kernel:**  $K_{ij} = \langle \phi(x^i), \phi(x^j) \rangle = k(x^i, x^j)$

**Classify New:**  $\operatorname{sign}(\sum a^i y^i k(x^i, x))$

**Polynomial Kernel t:**  $k(\langle x, x' \rangle + 1)^t$

**Gaussian Basis:**  $\exp(-\|x - x'\|_2^2)$

SVM (Linear Kernel):  $O(nd)$  train time,  $O(d)$  test time.

General Kernel:  $O(n^2 d)$  train time  $O(nd)$  test time.

## Decision Trees

Can do classification or regression, handle non-linear functions.

May fail on linear functions.

Start at one node, and split each. Select the pure node.

**Loss:**  $t^* = \operatorname{argmin}_t l(\{(x_i, y_i) : x_i \leq t\}) + l(\{(x_i, y_i) : x_i > t\})$

Let  $p_c = \text{frac of } S \text{ with label } c$ .  $\hat{y} = \operatorname{argmax}_c p_c$

**Misclassification Loss:**  $l(s) = 1 - p_y$

**Entropy Loss:**  $l(s) = - \sum_{\text{classes } c} p_c \log p_c$

**Gini Index Loss:**  $l(s) = \sum_{\text{classes } c} p_c (1 - p_c)$

**Regression:**  $l(s) = \min_p \sum_{i \in S} (y_i - p)^2 = \sum_{i \in S} (y_i - \bar{y}_s)^2$

We can stop based on run time, depth or splits.

Once a tree is fully grown we can prune it.

## Bagging

Training on empirical mean gives a variance of:  $E[\hat{\mu}] = \mu$

$$Var[\hat{\mu}] = Var[\frac{1}{n} \sum X_i] = \frac{1}{n^2} Var[\sum X_i] = \sigma^2/n$$

We can reduce variance by taking a sample of B points:

$$Var[\hat{\mu}] = Var[\frac{1}{B} \sum X_i] = \frac{1}{B^2} Var[\sum X_i] = \sigma^2/Bn$$

We can sample these points with replacement, and in practice this will work.

We aggregate by doing regression  $f(x) = \frac{1}{B} \sum f^j(x)$ .

Classification done by majority vote.

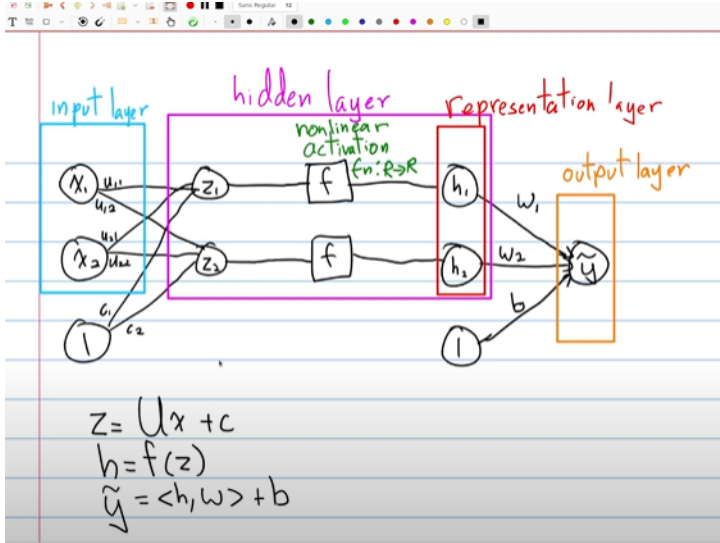
**Random Forests** Bootstrap but select  $\sqrt{d}$  features.

## Multilayer Perceptron

Neural Networks learn mapping from the data.

**Sigmoid:**  $\sigma(t) = \frac{1}{1+e^{-t}}$

$$\hat{y} = \frac{1}{1+\exp(-\langle w, x \rangle - b)}$$



**ReLU(t)** =  $\max(0, t)$

**Loss**  $l_\theta(x, y) = -\sum^m y_i \log y_i$

**Tanh(t)** =  $\frac{e^t - e^{-t}}{e^t + e^{-t}}$

**Gradient Descent**  $\theta^t = \theta^{t-1} - \eta \Delta L_{\theta^{t-1}}$

**Chain rule:**  $\frac{dz}{dx} = \frac{dz}{dy} \frac{dy}{dx}$

Any continuous function can be approximated well by a 2 layer nn.

## Multilayer Perceptron

Most neural networks have many parameters. We can minimize overfitting with:

Regularization loss, gradient descent, and equivalently:

$$\theta_t \leftarrow (1 - \eta \lambda) \theta_{t-1} - \eta \Delta L_{\theta_t - 1}(x, y)$$

We can drop nodes and use normalization of features:

$$\text{mean} = \frac{1}{n} \sum X_i, X_i \leftarrow X_i - \mu, \sigma_j^2 = \frac{1}{n} \sum X_{i,j}^2, X_{i,j} = X_{i,j} / \sigma_j$$

We do normalization on each batch, such that:

$$z^i = W^i h^{i-1} + b^i \text{ or } h^i = f(z^i)$$

We can do normalization on each neuron (batchnorm) or each layer.

**Batch GD:**  $\theta \leftarrow \theta - \eta * \frac{1}{n} \sum \delta l_\theta(x_i, y_i)$ , optimize gradient.

**Momentum:**  $v_t = \gamma v_{t-1} + (1 - \gamma) \mu, \theta_t \leftarrow \theta_{t-1} - v_t$  **RMSPprop** let  $g \in R^p, G_{t,i} = \sum^t g_{j,i}^2$  and  $\theta_t \leftarrow \theta_{t-1} - \frac{\mu}{\sqrt{G_{t,i} + \epsilon}} g_{t,i}$

- $\beta_1, \beta_2, \epsilon$  hyperparameters
  - $\beta_1 = 0.9, \beta_2 = 0.999, \epsilon = 10^{-8}$
- $m_{t,i} = \beta_1 m_{t-1,i} + (1 - \beta_1) g_{t,i}$  (momentum)
- $v_{t,i} = \beta_2 v_{t-1,i} + (1 - \beta_2) g_{t,i}^2$  (RMSPprop)
- $\hat{m}_{t,i} = \frac{m_{t,i}}{1 - \beta_1^t}, \hat{v}_{t,i} = \frac{v_{t,i}}{1 - \beta_2^t}$
- $\theta_{t,i} \leftarrow \theta_{t-1,i} - \frac{\eta}{\sqrt{\hat{v}_{t,i} + \epsilon}} \hat{m}_{t,i}$